


ANOVA



Statistical Experimental Design

Kris Sankaran | UW Madison | 21 September 2021

Feedback Response


- Restructuring lectures: Separating conceptual and code elements
- R Code: Only including essential R code
 - Have started an "R Cheatsheet"
- Pace / Volume: I will try to enunciate clearly, but please do raise your hand if I should slow down or repeat!
- R Review Session: 9/27 from 6:30 – 7:30pm
 - Hybrid: 1210 Medical Sciences Center and Zoom
- We will drop your lowest HW score

Today



- Book Sections: 3.1 – 3.4
- Online Notes: Week 3 [1] and [2]


Motivation



- ANOVA helps gauge the effects of ≥ 3 different treatments on a continuous response
 - How does the etch rate of a tool depend on its power?
 - How do different foods affect blood sugar?
 - How do several job training programs compare?
- It is an extension of two sample testing when there are 3 or more levels




Motivation



 Articles About 5,290,000 results (0.08 sec)



 Articles About 4,670,000 results (0.03 sec)


Model

ANOVA assumes,

$$y_{ij} = \mu + \tau_i + \epsilon_{ij}$$

where $i = 1, \dots, a$ and
 $j = 1, \dots, n$ and the errors
 $\epsilon_{ij} \sim \mathcal{N}(0, \sigma^2)$ are independent.

- i indexes different groups
- j indexes the samples within groups
- $N = na$ is the total number of samples



Hypothesis Testing

| Is there at least one group that differs from the rest?

Formally,

$$H_0 : \tau_1 = \cdots = \tau_a = 0$$

$$H_1 : \tau_i \neq 0 \text{ for at least one } i.$$

Q: What would be the version of the picture on the previous slide, if the null hypothesis were true?


ANOVA Identity

Before designing a test statistic, it helps to observe this identity,

$$\sum_{ij} (y_{ij} - \bar{y})^2 = n \sum_i (\bar{y}_i - \bar{y})^2 + \sum_{i,j} (y_{ij} - \bar{y}_i)^2$$

which is usually abbreviated as

$$SS_{\text{total}} = SS_{\text{treatment}} + SS_E.$$



Test Statistic

- If any of the groups are different from the global mean, then we expect $SS_{\text{treatment}}$ to be large
- How large is large enough?
- Consider the statistic,

$$\frac{MS_{\text{treatment}}}{MS_E}$$


where we define,

$$MS_{\text{treatment}} := \frac{1}{a - 1} SS_{\text{treatment}}$$

$$MS_E := \frac{1}{N - a} SS_E$$

Test Statistic

- It is not obvious, but reference distribution for this statistic is $F(a - 1, N - a)$.
- If this statistic is at a large quantile of that distribution, we conclude $\tau_i \neq 0$ for at least one i



Model Checking

We assumed,

$$y_{ij} = \mu + \tau_i + \epsilon_{ij}$$

with i.i.d. errors $\epsilon_{ij} \sim \mathcal{N}(0, \sigma^2)$


What could fail?

- The errors might not be normally distributed
- The variance might not be the same in each group
- The errors might not be independent
- There might be systematic variations besides the group deviations τ_i .

Residuals


Residuals are our best guess of what the true random error ϵ_{ij} , and so are useful for model checking,

$$\begin{aligned} e_{ij} &= y_{ij} - \hat{y}_{ij} \\ &= y_{ij} - (\hat{\mu} + \hat{\tau}_i) \end{aligned}$$




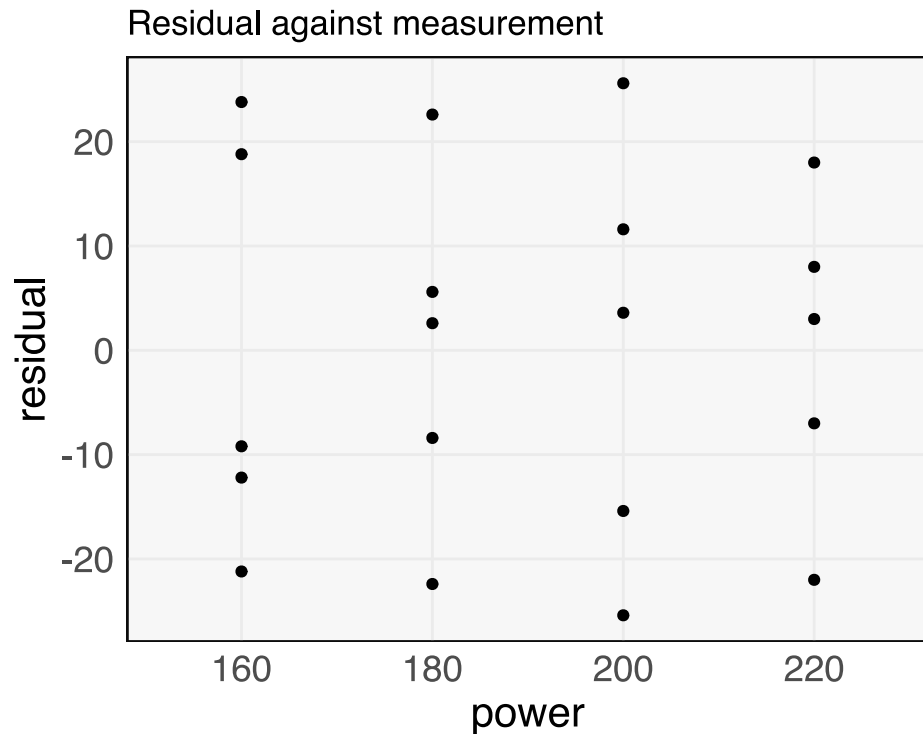
Normality assumption?

We can't check normality of ϵ_{ij} directly, but we can check normality of the residuals e_{ij} .





Equal variance across groups?

- Plot residuals against measured variables
- Plot residuals against fitted means



Unmeasured variation?

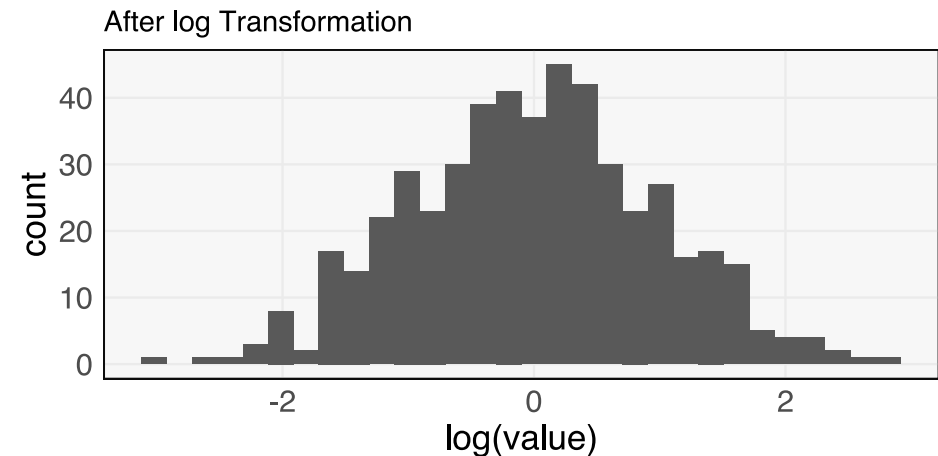
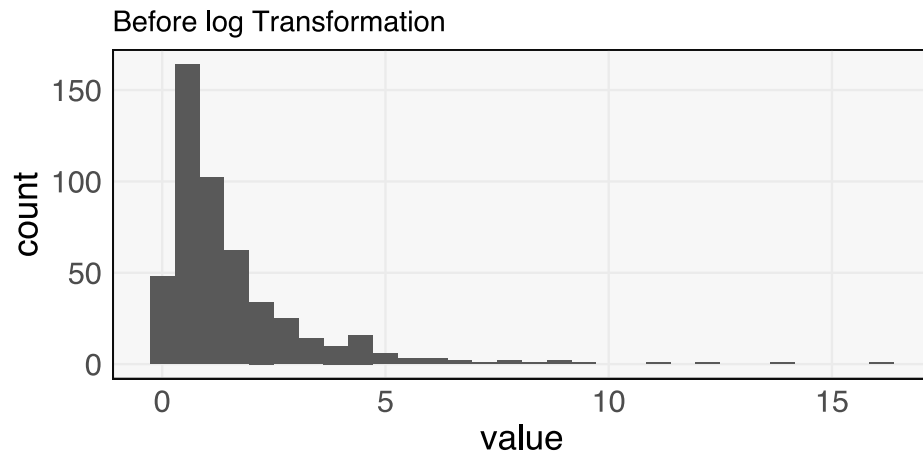
- Plot residuals against measured variables
- Plot residuals against fitted means



What if a check fails?

Transform the response variable to make it more bell-shaped

- Skewed, nonnegative data: Use $\log(x)$ or $\log(1 + x)$
- Count data: \sqrt{x} , $\sqrt{1 + x}$
- Proportions: $\arcsin(\sqrt{x})$



Code Implementation

Etch Rate Dataset


```
library(readr)
etch_rate <- read_csv("https://uwmadison.box.com/shared/static/vw3ldbgvg7rupt4tz3ditl1mpupw")
etch_rate$power <- as.factor(etch_rate$power) # want to think of power as distinct groups
head(etch_rate, 10)
```

```
## # A tibble: 10 × 3
##   power replicate  rate
##   <fct> <chr>      <dbl>
## 1 160    0b1         575
## 2 160    0b2         542
## 3 160    0b3         530
## 4 160    0b4         539
## 5 160    0b5         570
## 6 180    0b1         565
## 7 180    0b2         593
## 8 180    0b3         590
## 9 180    0b4         579
## 10 180    0b5         610
```

Plot the Data

- `ggplot()` expects a data.frame with the whole dataset
- `geom_point()` asks which columns to use for the x and y axis.

```
library(ggplot2)
ggplot(etch_rate) +
  geom_point(aes(power, rate))
```



ANOVA Hypothesis Test

- Which element in the table corresponds to $SS_{\text{treatment}}$?
- Which element in the table corresponds to $MS_{\text{treatment}}$?


```
fit <- lm(rate ~ power, data = etch_rate)
summary(aov(fit))
```

```
##              Df Sum Sq Mean Sq F value    Pr(>F)
## power          3  66871   22290     66.8 2.88e-09 ***
## Residuals     16   5339     334
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Check normality of residuals

- Use the `qqnorm` and `qqline` pair to make a QQ Plot


```
qqnorm(resid(fit))  
qqline(resid(fit), col = "red")
```



Plot residuals against factor

- First add the residual to the `data.frame`
- Then make the y -axis the residual


```
etch_rate$residual <- resid(fit)
ggplot(etch_rate) +
  geom_point(aes(power, residual))
```



Plot residuals against predicted value

- First add the fitted value to the `data.frame`
- Then make the x -axis the fitted value

```
etch_rate$predicted <- predict(fit)
ggplot(etch_rate) +
  geom_point(aes(predicted, residual))
```



Hint: Reshaping data

What if the etch dataset had been sent to us like this?

```
## # A tibble: 4 × 6
##   Power    Ob1    Ob2    Ob3    Ob4    Ob5
##   <fct> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 160      575    542    530    539    570
## 2 180      565    593    590    579    610
## 3 200      600    651    610    637    629
## 4 220      725    700    715    685    710
```

The `lm` function expects the outcome of interest to be in a single column.

Hint: Reshaping data

To reorganize the data into an acceptable form, we can use the `pivot_longer` function from the `tidyr` package.

```
library(tidyr)
etch <- pivot_longer(etch, -Power, names_to = "replicate", values_to = "etch_rate")
head(etch)
```

```
## # A tibble: 6 × 3
##   Power replicate etch_rate
##   <fct> <chr>      <dbl>
## 1 160    0b1         575
## 2 160    0b2         542
## 3 160    0b3         530
## 4 160    0b4         539
## 5 160    0b5         570
## 6 180    0b1         565
```

Exercise

This walks through problem 3.29 in the book.

A semiconductor manufacturer has developed three different methods for reducing particle counts on wafers. All three methods are tested on five different wafers and the after treatment particle count is obtained.

```
particles <- data.frame(  
  method = c("1", "2", "3"),  
  rep1 = c(31, 62, 53),  
  rep2 = c(10, 40, 27),  
  rep3 = c(21, 24, 120),  
  rep4 = c(4, 30, 97),  
  rep5 = c(1, 35, 68)  
)
```

particles

##	method	rep1	rep2	rep3	rep4	rep5
## 1	1	31	10	21	4	1
## 2	2	62	40	24	30	35
## 3	3	53	27	120	97	68

Exercise

- (1) Use `pivot_longer(particles, -method, ...)` to reshape the data so that the outcome is in a single column.
- (2) Use `lm` and `aov` to make an ANOVA table. Does the method detect any difference in the means across the three groups?
- (3) Plot the residuals against the group. Are the assumptions satisfied?
- (4) Apply a transformation to the response, fit an ANOVA model, and recheck normality of the residuals with a QQ plot. Stop when you are satisfied that the assumptions are met.

Solution

(1) `pivot_longer` can reshape the data.

```
library(tidyr)
particles <- pivot_longer(particles, -method, names_to = "replicate")
head(particles)
```

```
## # A tibble: 6 × 3
##   method replicate value
##   <chr>   <chr>     <dbl>
## 1 1      rep1        31
## 2 1      rep2        10
## 3 1      rep3        21
## 4 1      rep4         4
## 5 1      rep5         1
## 6 2      rep1        62
```

Solution

(2) The $MS_{\text{treatment}}$ is much larger than MS_E (4481.9 vs. 566.3), and the F test indicates a difference between methods at the 0.01 level.


```
fit <- lm(value ~ method, data = particles)
summary(aov(fit))
```

```
##              Df Sum Sq Mean Sq F value    Pr(>F)
## method         2   8964    4482   7.914 0.00643 **
## Residuals     12   6796     566
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Solution

(3) There appears to be very nonconstant variance across groups.


```
particles$residual <- resid(fit)
ggplot(particles) +
  geom_point(aes(method, residual))
```



Solution

(4) If we use a \sqrt{x} transformation, the difference between groups seems to be somewhat reduced, though not completely removed.

```
fit <- lm(sqrt(value) ~ method, data = particles)
particles$residual <- resid(fit)
ggplot(particles) +
  geom_point(aes(method, residual))
```



Solution

(4) Nonetheless, there still appears to be a significant difference between the groups' means.

```
##              Df Sum Sq Mean Sq F value    Pr(>F)
## method         2  63.90   31.95     9.84 0.00295 **
## Residuals     12  38.96    3.25
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```